

Software Engineering

Lecture 10 – Error Handling & Debugging

© 2015-19 Dr. Florian Echtler
Bauhaus-Universität Weimar
<florian.echtler@uni-weimar.de>

Error handling with Exceptions

- Goal: better structure for error handling
- Examples in C(++), but applicable to most other languages
- Exceptions are *thrown* (error reporting) and *caught* (error handling)

Error handling in C

- somewhat messy, requires lots of clean-up code
- errors often indicated through return value
- NULL for pointers
- value < 0 for int, often stored in global variable `errno`

```
int process_files() {  
  
    input1 = fopen("file1", "r");  
    if (input1 == NULL)  
        return errno;  
  
    input2 = fopen("file2", "r");  
    if (input2 == NULL) {  
        fclose(input1);  
        return errno;  
    }  
  
    input3 = fopen("file3", "r");  
    if (input3 == NULL) {  
        fclose(input1);  
        fclose(input2);  
        return errno;  
    }  
  
    // process files  
    [ ... ]  
    return 0;  
}
```

Error handling in C (2)

- Alternative method using goto
- “Goto considered harmful” (E. Dijkstra, 1968)
- Code using goto generally harder to read/understand
- Possible exception as shown in example

```
int process_files() {  
  
    input1 = fopen("file1", "r");  
    if (input1 == NULL)  
        goto error_1;  
  
    input2 = fopen("file2", "r");  
    if (input2 == NULL)  
        goto error_2;  
  
    input3 = fopen("file3", "r");  
    if (input3 == NULL)  
        goto error_3;  
  
    // process files  
    [ ... ]  
    return 0;  
  
error_3: fclose(input2);  
error_2: fclose(input1);  
error_1: return errno;  
}
```

Error handling in C++

- Local objects destroyed at end of scope
- No need for explicit clean-up if properly implemented in destructor
- Difficult with non-local objects

```
int process_files() {  
  
    MyFile input1("file1", "r");  
    if (!input1.valid())  
        return -1;  
  
    MyFile input2("file2", "r");  
    if (!input2.valid())  
        return -1;  
  
    MyFile input3("file3", "r");  
    If (!input3.valid())  
        return -1;  
  
    // process files  
    [ ... ]  
    return 0;  
}
```

Error handling in C++ (2)

- Solution: try/catch blocks for exceptions
- Classes from `std::` throw derivatives of `std::exception`
- `exception.what()` provides short text description of error

```
void process_files() {  
    try {  
        std::ifstream input1("file1");  
        std::ifstream input2("file2");  
        std::ifstream input3("file3");  
  
        // process files  
        [ ... ]  
    } catch (std::exception& e) {  
        // display error message  
        std::cout << e.what() << "\n";  
    }  
}
```

Exceptions with `throw/try/catch`

- Multiple catch blocks for different types are possible
- Rule (C++): “Throw by value, catch by reference”
- Java: Throw anything derived from `java.lang.Throwable` (e.g. subclass of `RuntimeException`)

```
try {  
    if (index < 0)  
        throw std::runtime_error(  
            "Index out of range" );  
  
    if (index > 10000)  
        throw -1;  
}  
  
// catch plain integer  
catch (int e) {  
    std::cout << "Error #";  
    std::cout << e << "\n";  
}  
  
// catch base class  
// of runtime_error  
catch (std::exception& e) {  
    std::cout << e.what() << "\n";  
}
```

Exception propagation

- Exceptions are propagated upwards until they are caught
- Use “catch-all” clauses for any type that does not have a specific handler
- Java catch-all:
catch (Exception ex)

```
void throw_something() {  
    // no try...catch here  
    throw std::runtime_error();  
}  
  
void caller() {  
  
    try {  
        throw_something();  
    }  
  
    // exception caught here  
    catch (std::exception& e) {  
        std::cout << e.what() << "\n";  
    }  
  
    // “catch-all” for any type  
    catch (...) {  
        std::cout << “default  
        exception handler\n”;  
    }  
}
```


Debugging

- The “basic” approach: `printf()` & friends
 - Not very flexible, modifiable at compile-time only
 - Mental “reverse-engineering” of program state
 - Can slow program down (sometimes massively)
 - Advantages?
- The “smart” approach: debugger
 - Stop program at arbitrary point
 - Advance execution step-by-step
 - View arbitrary variables/objects

Debugger commands

Exact command names differ between debuggers (gdb, VS, XCode, ...), functionality remains same

- Set Breakpoint
 - Execution stops at specified line or method
- Set Watchpoint
 - Execution stops when variable is changed

Debugger commands (2)

- Step → Execute a single line of code
 - Step Into: enter any called functions
 - Step Over: call functions as a whole and stop
- Inspect/Watch Variable
 - Value of variable is shown while stepping
- Run/Continue
 - Execute until breakpoint/watchpoint hit

Stack frames

- Stack: dedicated memory for each thread
 - keeps track of called functions and their local vars
 - grows *downward* as opposed to heap (malloc, new)
- Stack frame: local variables & return address of currently executed function
- Debugger can switch to earlier stack frames, i.e. inspect caller's context
- List of stack frames = *call stack*

Call stack

- Minimal example with up to 3 stack frames, state when execution paused in line 2

→

```
1 int main() {  
2     printf("Hello!");  
3     func1();  
4     printf("Goodbye!");  
5 }
```

```
6 void func1() {  
7     int i = 1;  
8     func2(i);  
9     i = 42;  
10 }
```

```
11 void func2(int a) {  
12     int b = a;  
13 }
```

↓ Start of stack:
Frame for main()

Call stack (2)

- Minimal example with up to 3 stack frames, state when execution paused in line 7

```
1 int main() {  
2     printf("Hello!");  
3     func1();  
4     printf("Goodbye!");  
5 }
```

→

```
6 void func1() {  
7     int i = 1;  
8     func2(i);  
9     i = 42;  
10 }
```

```
11 void func2(int a) {  
12     int b = a;  
13 }
```

Start of stack:

Frame for main()

Frame for func1()

Return to main, line 4

Storage for int i

Call stack (3)

- Minimal example with up to 3 stack frames, state when execution paused in line 12

```

1 int main() {
2     printf("Hello!");
3     func1();
4     printf("Goodbye!");
5 }

```

```

6 void func1() {
7     int i = 1;
8     func2(i);
9     i = 42;
10 }

```

```

11 void func2(int a) {
12     int b = a;
13 }

```

Start of stack:

Frame for main()

Frame for func1()

Return to main, line 4

Storage for int i

Frame for func2()

Return to func1, line 9

Storage for int b

Storage for int a


Debugging information

- Debugger does not know names of functions, variables etc.
→ can only show memory addrs
- “Human-readable” names must be compiled into executable,
e.g. using -g flag (gcc/javac)
- Binary grows larger, sometimes slower
→ usually not enabled for release builds

Buffer Overflow (1)


Source (FU): <http://arstechnica.com/security/...-the-buffer-overflow/>

```
void get_name() {
    char name[64];
    printf("Your name? ");
    std::gets(name);
    printf("Hi, %s.", name);
}
```



Source (FU): <http://arstechnica.com/security/...-the-buffer-overflow/>

- ```
top of stack
esp = 0x0019fee0
```






# Buffer Overflow (4)

Source (FU): <http://arstechnica.com/security/...-the-buffer-overflow/>

- Return address can be overwritten with pointer to buffer itself
- Possible to insert *and execute* arbitrary code!
- Root cause of many historic security issues (e.g. Morris worm)

top of stack  
esp = 0x0019fee0



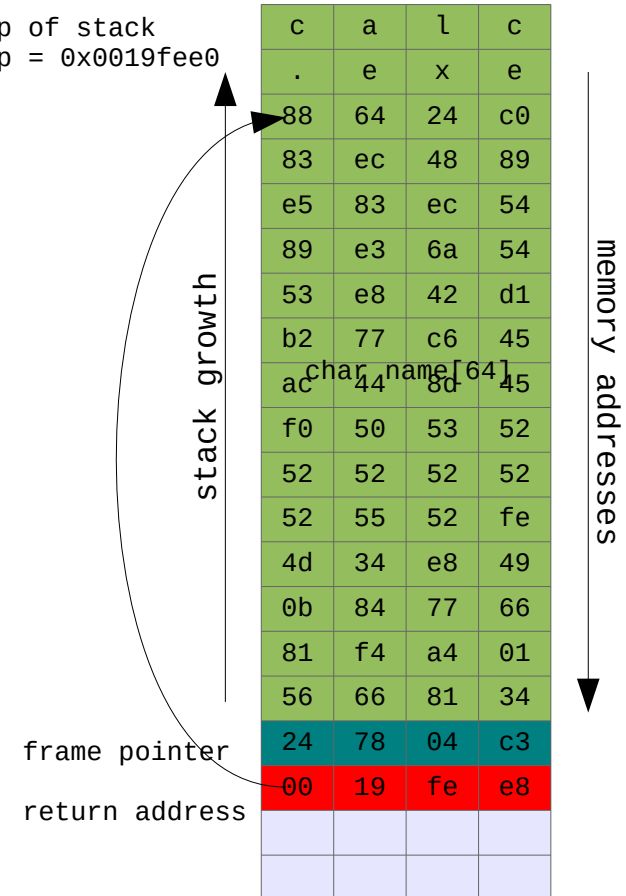


# Buffer Overflow (5)

Source (FU): <http://arstechnica.com/security/...-the-buffer-overflow/>

- Modern OS contain multiple countermeasures
- Possible mitigation strategies?
  - Address layout randomization
  - Stack “canary” values
  - No-Execute-Bit (NX)

top of stack  
esp = 0x0019fee0



# Questions/Comments?

